

# Computer Vision Systems Programming VO

## A Recap of Image Processing

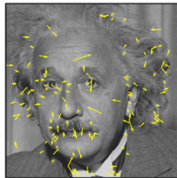
Christopher Pramerdorfer

Computer Vision Lab, Vienna University of Technology

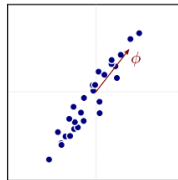
# Topics

## A brief recap of Image Processing (IP)

- ▶ Assuming you are already familiar with IP
- ▶ Focus on methods that are widely used in practice



Images from Prince 2012



# Relation of IP and CV

IP encompasses operations that

- ▶ Take images as input
- ▶ Produce images or representations (e.g. descriptors)

We regard IP as preprocessing for CV

- ▶ IP has great influence on CV performance

# Relation of IP and CV

CV is all about

- ▶ Inferring some world state  $w$
- ▶ From measurements  $x$

We use IP to obtain a suitable  $x$  from images

Suitable often means

- ▶ Distinctive features
- ▶ That are robust and covariant

# Contrast Normalization

Reduce variation due to contrast and intensity changes

We cover two techniques

- ▶ Whitening
- ▶ Histogram equalization



Images from Prince 2012

# Contrast Normalization

## Whitening

Transform pixel values so that

- ▶ Their mean is zero
- ▶ Their variance is one



Images from Prince 2012

# Contrast Normalization

## Whitening – Matlab Implementation

```
img = single(rgb2gray(imread('image.png'))); % load  
m = mean(img(:)); % compute mean  
s = std(img(:)); % compute standard deviation  
whitened = (img - m) / s; % normalize
```

# Contrast Normalization

## Histogram equalization

Transform pixel values so that distribution is “flat”

- ▶ Cumulative histogram linear over value range



Images from Prince 2012



# Contrast Normalization

## Histogram Equalization – C++ Implementation

```
// cv = OpenCV namespace  
cv::Mat img, equalized; // storage  
img = cv::imread("image.png", cv::IMREAD_GRAYSCALE); // load  
cv::equalizeHist(img, equalized); // normalize
```

# Noise Reduction and Change Detection

Reduce image noise

Locate intensity changes

Often accomplished via **linear filtering**

- ▶ Pixel values linear combination of neighbor values
- ▶ Computed via convolution (or correlation)

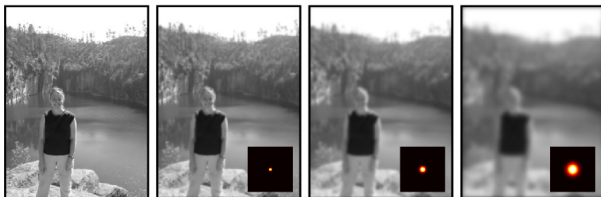
$$f'(x, y) = \sum_{i,j} f(x - i, y - j)h(i, j)$$

# Noise Reduction and Change Detection

## Noise Reduction via Blurring

Use a 2D Gaussian as kernel  $h$ :

$$h(i, j) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{i^2 + j^2}{2\sigma^2}\right)$$



Images from Prince 2012

# Noise Reduction and Change Detection

## Noise Reduction via Blurring – Matlab Implementation

```
img = rgb2gray(imread('image.png')); % load  
h = fspecial('gaussian', [3 3], 0.5); % create Gaussian kernel  
filtered = imfilter(img, h); % filter
```

# Noise Reduction and Change Detection

## Change Detection via LoG Filtering

Use a Laplacian of Gaussian (LoG) filter as kernel  $h$

- ▶ Gaussian for noise reduction
- ▶ Laplacian approximates  $\nabla^2 = f_{xx} + f_{yy}$

LoG filters respond to intensity changes

- ▶ Regardless of direction
- ▶ At a frequency defined by  $\sigma$  of Gaussian

Substrate for SIFT interest points

# Noise Reduction and Change Detection

## Change Detection via LoG Filtering



Images from Prince 2012

# Noise Reduction and Change Detection

## Change Detection via LoG Filtering – Matlab Implementation

```
img = rgb2gray(imread('image.png')); % load  
h = fspecial('log', [3 3], 0.5); % create LoG kernel  
filtered = imfilter(img, h); % filter
```

# Noise Reduction and Change Detection

## Change Detection via Gabor Filtering

Use a Gabor filter as kernel  $h$ , which consists of

- ▶ A Gaussian for noise reduction
- ▶ A Sinusoid for change detection

Gabor filters respond to intensity changes at a

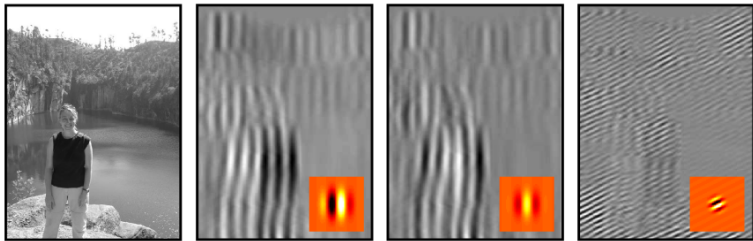
- ▶ Phase and orientation defined by the Sinusoid
- ▶ Frequency defined by the Gaussian and Sinusoid

Substrate for object recognition and scene understanding



# Noise Reduction and Change Detection

## Change Detection via Gabor Filtering



Images from Prince 2012

# Noise Reduction and Change Detection

## Change Detection via Gabor Filtering – C++ Implementation

```
cv::Mat img, gabor; // storage  
img = cv::imread("image.png", cv::IMREAD_GRAYSCALE); // load  
h = cv::getGaborKernel(...); // create Gabor kernel  
cv::filter2D(img, gabor, CV_32F, h); // filter
```

# Interest Point Detection

Interest points are image features that

- ▶ Are local (precise location, little spatial extent)
- ▶ Can be detected reliably in multiple images of same object

Which implies that they are

- ▶ Invariant to image transformations
- ▶ Robust with respect to noise

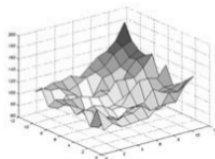
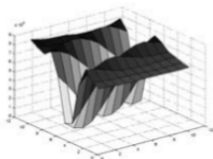
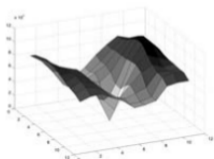
# Interest Point Detection

## Harris Corner Detector

Corners characterized by intensity change in multiple directions

Harris corner detector exploits this by

- ▶ Checking gradient distribution in local neighborhood
- ▶ Corner: gradient distribution has two large eigenvalues



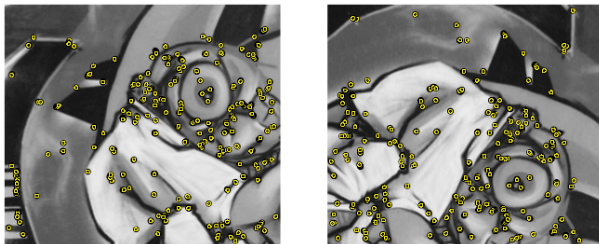
Images from Szeliski 2010

# Interest Point Detection

## Harris Corner Detector

### Harris interest points

- ▶ Are invariant to translation and rotation
- ▶ Stable under varying lighting conditions



Images from Tuytelaars and Mikolajczyk 2008

# Interest Point Detection

## Harris Corner Detector – Matlab Implementation

```
img = rgb2gray(imread('image.png')); % load  
corners = corner(img, 'Harris', maxNum); % detect corners
```

# Interest Point Detection

## SIFT Detector

Scale invariant blob detector

- ▶ A blob is an image region with similar intensity

Blob detection accomplished via LoG filtering

- ▶ LoG filter responds to blobs of size that depends on  $\sigma$


Scale invariance is achieved by

- ▶ Applying LoG filter with multiple  $\sigma$
- ▶ Finding local maxima in resulting scale-space


Repeated LoG approximated by Differences of Gaussians (DoGs)

# Interest Point Detection

## SIFT Detector – Scale Space



Images from Prince 2012





# Interest Point Detection

## SIFT Detector

Local maxima are

- ▶ Localized to sub-voxel accuracy
- ▶ Discarded unless on corners
- ▶ Assigned an orientation via gradient histograms

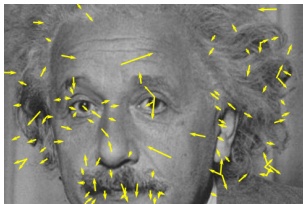


Image from Prince 2012

# Interest Point Detection

## SIFT Detector – C++ Implementation

```
cv::Mat img = cv::imread("image.png", cv::IMREAD_GRAYSCALE);  
cv::SiftFeatureDetector det(20, 10); // create detector  
std::vector<cv::KeyPoint> kps; // keypoint storage  
det.detect(img, kps); // detect keypoints
```

# Local Descriptors

Compact representations of contents of an image region

Usually computed at interest point locations

Invariant in conjunction with suitable interest points

Pool information locally to achieve robustness

- ▶ Often accomplished via histograms

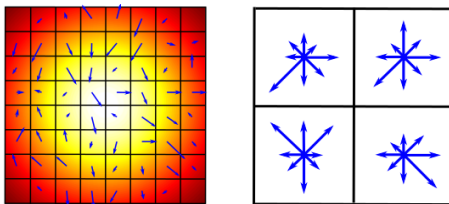
# Local Descriptors

## SIFT Descriptor

Computed from gradient histograms

Usually used together with SIFT interest points

- Compensate for scale, rotation



Images from Prince 2012

# Local Descriptors

## SIFT Descriptor

SIFT descriptors are

- ▶ Invariant to scale and rotation (interest points)
- ▶ Invariant to global intensity changes (gradients)
- ▶ Robust to small affine transformations (pooling)

# Local Descriptors

## SIFT Descriptor – C++ Implementation

```
// img and kps from previous example  
cv::Mat descriptors; // storage (n by 128)  
cv::SiftDescriptorExtractor ex; // create extractor  
ex.compute(img, kps, descriptors); // compute descriptors
```

# Dimensionality Reduction

Sometimes desirable to reduce the dimensionality of  $\mathbf{x}$

- ▶ Makes learning and inference more efficient
- ▶ Can improve generalization performance
- ▶ Facilitates data visualization

Goal is to find transformation from  $\mathbf{x}$  to  $\mathbf{v}$

- ▶ With  $v = \dim(\mathbf{v}) < x = \dim(\mathbf{x})$
- ▶ That minimizes the information loss

# Dimensionality Reduction

## Principal Component Analysis

Principal Component Analysis (PCA) yields

- ▶ The orthogonal transformation  $\phi$  to a  $v$ -dimensional subspace
- ▶ That minimizes  $\sum_i \|\mathbf{x}_i - \phi^{-1} \mathbf{v}_i\|_2$  (assuming zero mean)

Accomplished if  $\phi = [\phi_1, \dots, \phi_v]_{v \times x}^T$

- ▶ With  $\phi_k$  being the  $k$ th largest eigenvectors of  $\mathbf{X}\mathbf{X}^T$
- ▶ Where  $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_I]$

Projects  $\mathbf{x}_i$  onto hyperplane

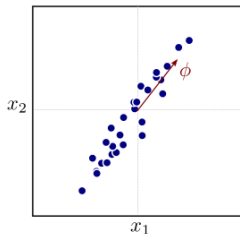
- ▶ Spanned by  $v$  largest axes of covariance ellipsoid

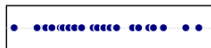


# Dimensionality Reduction

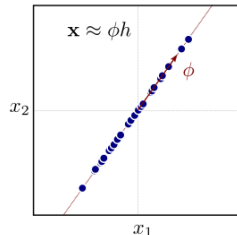
## Principal Component Analysis – Example

In this case  $h = v$  and  $\phi = \phi^{-1} = \phi^T$



$$h = \phi^T \mathbf{x}$$


A 1D plot showing the principal component  $h$  (blue dots) along a horizontal axis. The points are tightly clustered along a single line, indicating that the data is effectively 1D when projected onto the principal component.



Images from Prince 2012

# Dimensionality Reduction

## Principal Component Analysis – Remarks

PCA is dependent on the units of  $x_i$

- ▶ So standardize (whiten) your data

PCA is unsupervised and linear

- ▶ Better methods for supervised case (e.g. LDA)
- ▶ More powerful non-linear methods (e.g. autoencoders)
- ▶ But generally superfluous with modern learning methods

# Dimensionality Reduction

## Principal Component Analysis – Matlab Implementation

```
pc = princomp(X'); % compute all principal components (sorted)
phi = c(:,1:v)'; % keep only v "largest" ones and transpose
V = phi * X; % project to lower dimension
Xe = phi' * V % reconstruct X
```

# Summary

We utilize IP to obtain  $\mathbf{x}$  from an image

- ▶ That is for feature extraction

We want  $\mathbf{x}$  to be distinctive, invariant, robust, concise

- ▶ And we have seen methods to achieve this

We have only scratched the surface of IP

- ▶ See literature

# Bibliography

- Prince, S.J.D. (2012). **Computer Vision: Models Learning and Inference**. Cambridge University Press.
- Szeliski, Richard (2010). **Computer vision: algorithms and applications**. Springer.
- Tuytelaars, Tinne and Krystian Mikolajczyk (2008). **Local invariant feature detectors: a survey**. Foundations and Trends in Computer Graphics and Vision.