

# Computer Vision Systems Programming VO

## Models vs. Algorithms

Christopher Pramerdorfer


Computer Vision Lab, Vienna University of Technology

# Topics

Models vs. algorithms

Approaching computer vision problems

Numerical optimization



# Solving CV Problems

CV is about inferring some world state  $\mathbf{w}$  from measurements  $\mathbf{x}$

$\mathbf{x}$  is a feature vector extracted from images

$\mathbf{w} \in \mathbb{R}^m$  : **regression** problem (e.g. stereo)

$\mathbf{w} \in \mathbb{N}^m$  : **classification** problem (e.g. face recognition)

# Solving CV Problems

## How Not To Do It

Don't think in terms of **algorithms**

- ▶ First I'll convert my image to grayscale
- ▶ Then I'll threshold it at some value
- ▶ Then I'll clean up holes via morphology

Such solutions

- ▶ Rely on many magic numbers that all interact
- ▶ Do not encode uncertainty about  $w$

# Solving CV Problems

## How To Do It

Think about how your data came into being

Use this information to **model** the relationship between  $x$  and  $w$

Ideally use probabilistic models as inference is ill-posed

“Say **what** you want to do, not **how** you are going to to it”

~ A. Fitzgibbon

# Solving CV Problems

## How To Do It

On this basis, CV is about

- ▶ Defining a **model** that relates  $\mathbf{x}$  and  $\mathbf{w}$
- ▶ **Learning** its parameters  $\theta$  from **training samples**
- ▶ Using an **inference algorithm** that returns  $\Pr(\mathbf{w}|\mathbf{x}, \theta)$

# Solving CV Problems

## Example: Background Subtraction

We want to segment images in foreground and background

- ▶ Based on the perceived brightness of objects
- ▶ In scenes with a constant background
- ▶ On a per-pixel basis



Image from Prince 2012

# Solving CV Problems

## Example: Background Subtraction

Therefore we have  $w \in \{0, 1\}$  (classification problem)

Given our problem, we model

$$\Pr(x|w = 0) = \text{Norm}_x(\mu, \sigma^2)$$

$$\Pr(x|w = 1) = \text{Uniform}_x(0, 255)$$

$$\Pr(w) = \text{Bern}_w(\lambda)$$



# Solving CV Problems

## Example: Background Subtraction

We assume  $\lambda = 0.5$

And learn  $\theta = (\mu, \sigma)$  from training samples  $\{x_i\}_{i=1}^n$

- ▶ Using the maximum likelihood method
- ▶ Assuming that the training samples are independent

We thus learn

$$\begin{aligned}\hat{\theta} &= \arg \max_{\theta} \Pr(x_1, \dots, x_n | \theta) \\ &= \arg \max_{\theta} \prod_{i=1}^n \Pr(x_i | \theta)\end{aligned}$$

# Solving CV Problems

## Example: Background Subtraction

Continuing from before

$$\begin{aligned}\hat{\theta} &= \arg \max_{\theta} \prod_{i=1}^n \text{Norm}_{x_i}(\mu, \sigma^2) \\ &= \arg \max_{\theta} \sum_{i=1}^n \log \text{Norm}_{x_i}(\mu, \sigma^2) \\ &= \arg \max_{\theta} \left( -0.5n \log \sigma^2 - 0.5n \log(2\pi) - 0.5 \sum_{i=1}^n \frac{(x_i - \mu)^2}{\sigma^2} \right)\end{aligned}$$

# Solving CV Problems

## Example: Background Subtraction

If we differentiate and equate to zero, we obtain

$$\begin{aligned}\hat{\mu} &= \frac{1}{n} \sum_{i=1}^n x_i \\ \hat{\sigma}^2 &= \frac{1}{n} \sum_{i=1}^n (x_i - \hat{\mu})^2\end{aligned}$$

We have derived **algorithms** for finding our model parameters

# Models vs. Algorithms

Models describe the relationship between  $\mathbf{x}$  and  $\mathbf{w}$

Algorithms are recipes for finding  $\theta$

Model	Algorithm
Gaussian	Mean & variance
Linear models*	Linear least squares
Structure from motion*	Bundle adjustment
Rigid motion between point sets*	ICP

\* Assuming zero-mean Gaussian noise

# Inference

In the example case we have modeled  $\Pr(\mathbf{x}|\mathbf{w})$

- ▶ Such models are called **generative**
- ▶ In this case, we perform inference using Bayes' rule

$$\Pr(\mathbf{w}|\mathbf{x}) = \frac{\Pr(\mathbf{x}|\mathbf{w}) \Pr(\mathbf{w})}{\Pr(\mathbf{x})}$$

We can also model  $\Pr(\mathbf{w}|\mathbf{x})$  directly

- ▶ Such models are called **discriminative**
- ▶ In this case, we can perform inference directly

# Inference

## Example: Background Subtraction

Assume  $x = 100, \mu = 95, \sigma = 2$

In this case, we obtain

- ▶  $\Pr(w = 0) = \Pr(w = 1) = 0.5$
- ▶  $\Pr(x = 100|w = 0) = \text{Norm}_{100}(90, 2^2) \approx 0.009$
- ▶  $\Pr(x = 100) = \sum_{i \in \{0,1\}} \Pr(x = 100|w = i) \Pr(w = i) \approx .007$
- ▶  $\Pr(w|x = 100) \approx (0.69, 0.31)$

Therefore we classify the pixel as background

# Complex Models




Image from <https://www.youtube.com/watch?v=sQegEro5Bfo>

# Optimization

Learning  $\theta$  is an **optimization problem**

- ▶ We seek  $\hat{\theta}$  that maximizes some objective function  $f(\theta)$
- ▶ E.g. before we sought  $\hat{\theta} = \arg \max_{\theta} \Pr(x_1, \dots, x_n | \theta)$

There are algorithms available for finding  $\theta$  of many models

- ▶ But this is not always the case

Therefore, a brief intro on optimization ...



# Optimization

We seek  $\hat{\theta} = \arg \max_{\theta} f(\theta) = \arg \min_{\theta} -f(\theta)$

- ▶ Sometimes  $\hat{\theta}$  can be found analytically (like before)
- ▶ Often we have to resort to iterative methods

Iterative methods

- ▶ Start with an estimate  $\theta^{[0]}$
- ▶ Improve this estimate iteratively using local information
- ▶ Until no more improvement can be made

# Optimization

## Iterative Methods




Image adapted from Prince 2012

# Optimization

## Iterative Methods

This means we need a good initial estimate  $\theta^{[0]}$

- ▶ Otherwise we might get stuck in a **local minimum**
- ▶ Unless  $f$  is **convex**




Image from Prince 2012

Iterative methods work by alternatively

- ▶ Choosing a search direction using local information
- ▶ Searching the minimum along that direction in some interval
- ▶ Updating  $\theta$  according to the found minimum

Local information from first and second order derivatives


Many different optimization algorithms

- ▶ Steepest decent, Newton, Levenberg-Marquardt, BFGS ...
- ▶ Some can be tested at <https://awful.codeplex.com/>

# Optimization

## Iterative Methods – Simple Example

We seek the minimum of  $f(\boldsymbol{\theta}) = (\theta_1 - 2)^2 + (\theta_2 - 1)^2$



Finding the minimum with Python and SciPy

```
def f(x):  
    return np.power(x[0]-2,2) + np.power(x[1]-1,2)  
  
scipy.optimize.minimize(f, [5, 5]) # returns [2, 1]
```

### Finding the minimum with Matlab

```
fminunc(@(x)(x(1)-2)^2 + (x(2)-1)^2, [5, 5]) % returns [2, 1]
```

# Summary

We have seen how (not) to approach CV problems

- ▶ See Prince 2012 for more information on CV models

Optimization is fundamental in vision problems

- ▶ We briefly discussed why, and how optimization works



Books for further information on these topics

- ▶ Prince 2012 (<http://computervisionmodels.com/>)
- ▶ Bishop et al. 2006
- ▶ Boyd and Vandenberghe 2009
- ▶ Wright and Nocedal 1999

# Bibliography

Bishop, Christopher et al. (2006). **Pattern Recognition and Machine Learning**. Springer.

Boyd, Stephen and Lieven Vandenberghe (2009). **Convex Optimization**. Cambridge University Press.

Prince, S.J.D. (2012). **Computer Vision: Models Learning and Inference**. Cambridge University Press.

Wright, SJ and J Nocedal (1999). **Numerical optimization**. Springer.