

Computer Vision Systems Programming VO

Object Category Recognition

Christopher Pramerdorfer

Computer Vision Lab, Vienna University of Technology

Topics

- Scene classification using the bag of words model
- Fast face detection using boosted Haar features
- Convolutional neural networks for large-scale problems




Image adaoted from Kavukcuoglu 2011

Scene Classification

We want to distinguish between c scene categories

- So $w \in \{0, \dots, c - 1\}$ (classification problem)

Street Scenes



Sea Scenes



Image adapted from Prince 2012

Scene Classification

Bag of Visual Words

We represent an image as a collection of **visual words**

- Images can be compared based on visual word distribution




Image from Prince 2012

Scene Classification

Bag of Visual Words

Visual words are learned from an image collection

- ▶ Compute (SIFT) keypoints and descriptors for all images
- ▶ Cluster descriptors into k clusters using k -means
- ▶ k cluster means represent visual words




Image from Grauman and Leibe 2011

Scene Classification

Bag of Visual Words

Visual word distribution $\mathbf{x} \in \mathbb{N}^k$ of image obtained by

- ▶ Computing keypoints and descriptors
- ▶ Assigning each feature to closest visual word
- ▶ Summing up the assignment counts for each visual word




Image from Grauman and Leibe 2011

Scene Classification

Bag of Visual Words

This image representation is called **bag of (visual) words**

Now that we have x we can select and learn a suitable model

- ▶ SVMs are often used in the literature
- ▶ For a probabilistic alternative see Prince 2012

Scene Classification

Bag of Visual Words – Remarks

Many improvements to this model exist

- ▶ Better clustering schemes
- ▶ Fuzzy assignment to visual words
- ▶ Spatial information (constellation model)

Popular and can work well, but no longer state of the art

Scene Classification

Bag of Visual Words Using OpenCV

```
// init SIFT
cv::Ptr<cv::FeatureDetector> kp = cv::FeatureDetector::create("SIFT");
cv::Ptr<cv::DescriptorExtractor> desc = cv::DescriptorExtractor::create("SIFT");

// compute visual words from training data
const int k = 50; // number of visual words
cv::BOWKMeansTrainer trainer(k);

for(const cv::Mat& im : images) { // std::vector of training images
    std::vector<cv::KeyPoint> keypoints; kp->detect(im, keypoints);
    cv::Mat descriptors; desc->compute(im, keypoints, descriptors);
    trainer.add(descriptors);
}

cv::Mat visualWords = trainer.cluster(); // k*128 (SIFT dimension)
```

Scene Classification

Bag of Visual Words Using OpenCV

```
// setup visual word frequency (our x) extractor
cv::Ptr<cv::DescriptorMatcher> fm = cv::makePtr<cv::BFMatcher>(cv::NORM_L2);
cv::BOWImgDescriptorExtractor extractor(desc, fm);
extractor.setVocabulary(visualWords);

// compute x for all training images
cv::Mat xTrain(images.size(), k, CV_32FC1);
for(std::size_t i = 0; i != images.size(); i++) {
    std::vector<cv::KeyPoint> keypoints; kp->detect(images[i], keypoints);
    cv::Mat x; extractor.compute(images[i], keypoints, x);
    xTrain(cv::Rect(0, i, k, 1)) = x;
}

// and corresponding w
cv::Mat wTrain(images.size(), 1, CV_32FC1); // fill me
```

Scene Classification

Bag of Visual Words Using OpenCV

```
// train our model (we use an SVM)
```

```
CvSVM svm;
```

```
svm.train(xTrain, wTrain);
```

```
// now we can predict the class of new images
```

```
std::vector<cv::KeyPoint> keypoints; kp->detect(newImage, keypoints);
```

```
cv::Mat x; extractor.compute(newImage, keypoints, x);
```

```
float w = svm.predict(x); // predicted class label
```


Face Detection




Image from olympus-europa.com

Face Detection

We don't know where the faces are so we

- ▶ Slide a fixed-size window over the image
- ▶ Compute $\Pr(w|\mathbf{x})$ for each window ($w = 1$ if face, 0 if not)



Face Detection

Selecting x

Which are good features for this task?

Must be fast to compute (many windows)

Must be robust to illumination, so we use gradient information

Different approaches to encoding gradient information

- ▶ Compute gradients, pool orientations in blocks (e.g. SIFT)
- ▶ Use a collection of Gabor filters

Face Detection

Selecting x

We want something faster

- ▶ So we use a “blocky” approximation of Gabor filters
- ▶ Difference between rectangular subwindows (**Haar features**)
- ▶ Can be computed in constant time using integral images




Image adapted from Prince 2012

Face Detection

Selecting \mathbf{x}

Computing a Haar feature yields a scalar f_i

We define $\mathbf{x} = (x_1, \dots, x_I)$ with $x_i = \text{heaviside}(f_i - t_i)$

I is very large (but finite)

- ▶ Different Haar features, subwindow locations, t_i
- ▶ We thus learn which features x_i work best

We model $\Pr(w|\mathbf{x})$ as a weighted sum of a feature subset

$$\Pr(w|\mathbf{x}) \propto a = \phi_0 + \sum_k \phi_k x_k \quad (1 \leq k \leq I)$$

And learn the parameters $\theta = (\phi_0, \phi_k, x_k)$ from training samples

- ▶ For each x_k in a large precomputed set
- ▶ Find optimal ϕ_0, ϕ_k
- ▶ Add best x_k to sum and repeat

This incremental approach is called **boosting**

Face Detection

Boosting

We stop adding features at some point

- ▶ If the classification error no longer decreases
- ▶ After a specified maximum number of iterations

We end up with K good features, $K \ll I$

- ▶ For prediction we compute only these K features
- ▶ We stop if $P(w = 0) > t$ after processing $J \ll K$ features

Face Detection

Boosting

If we don't care about probabilities we choose $w = \text{heaviside}(a)$

If we do we use **logistic regression**, $\Pr(w|\mathbf{x}) = \text{Bern}_w(\text{sig}(a))$

- ▶ We model w as a Bernoulli distribution
- ▶ Pass a through a logistic sigmoid to map it to $[0, 1]$
- ▶ Called **logitboost** in this context

Face Detection

Viola & Jones Face Detector

This method was proposed in Viola and Jones 2001

Very efficient, ideal for digital cameras

Trades off efficiency for accuracy

- ▶ Features capture gradients coarsely, no color information
- ▶ More powerful but slower methods exist

Not invariant to scale changes (fixed-size window)

- ▶ Repeat detection at different image scales

Face Detection

Viola & Jones Face Detector in OpenCV

Detect faces using a pretrained model

OpenCV also supports training

```
# detect faces using a pretrained cascade
```

```
image = cv2.imread('faces.jpg', cv2.IMREAD_GRAYSCALE)
```

```
cascade = cv2.CascadeClassifier('haarcascade_frontalface_default.xml')
```

```
faces = cascade.detectMultiScale(image) # should tune parameters
```


Selecting good features x for object recognition is challenging

- ▶ Why we previously learned x

Learned features were low-level

- ▶ Based on SIFT descriptors or Haar wavelets

We want task-specific high-level features

- ▶ Virtually impossible to design manually
- ▶ So we learn them as well

Deep Learning

We learn these features hierarchically

- ▶ Model consists of layers
- ▶ The higher up the layer, the higher-level the feature
- ▶ Features in layer n are based on those in layer $n - 1$

This results in a *deep* model, hence **deep learning**

At the same time we learn to predict w

Deep Learning




Image from Bengio, Goodfellow, and Courville 2015

Bibliography I

- Bengio, Yoshua, Ian Goodfellow, and Aaron Courville (2015). *Deep Learning (Draft)*. MIT Press.
- Grauman, Kristen and Bastian Leibe (2011). *Visual object recognition*. Morgan & Claypool.
- Kavukcuoglu, Koray (2011). *Learning feature hierarchies for object recognition*. PhD thesis.
- Prince, S.J.D. (2012). *Computer Vision: Models Learning and Inference*. Cambridge University Press.
- Viola, Paul and Michael Jones (2001). *Rapid object detection using a boosted cascade of simple features*.