

Computer Vision Systems Programming VO

Deep Learning

Christopher Pramerdorfer

Computer Vision Lab, Vienna University of Technology

Topics



Deep learning motivation

Multilayer perceptrons

Convolutional neural networks

Deep learning applications


80322-4129
40004
07879



Images from LeCun et al. 1989, Taigman et al. 2013, image-net.org

Object Recognition

Traditional Approach



Object Recognition

Traditional Approach

Problem: how to choose the representation/features?

“General” features not optimal

- ▶ Not tuned to task at hand, low-level

Designing task-specific features is complex

- ▶ Virtually impossible to do optimally

Object Recognition

Deep Learning

Solution: learn representation as well

Learning high-level representations directly is difficult

Deep Learning (DL) solves this

- ▶ By learning a hierarchy of representations
- ▶ Layers in hierarchy build upon each other

Object Recognition

Deep Learning

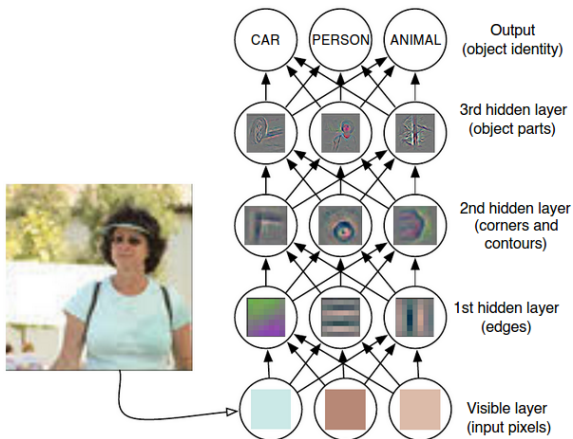


Image from Bengio, Goodfellow, and Courville 2014

Object Recognition

Deep Learning

n levels of features/representations

Learned jointly with the output mapping



Multilayer Perceptrons

DL is usually realized using MultiLayer Perceptrons (MLPs)




Image from Bishop 2006

Multilayer Perceptrons

The Perceptron

Binary linear classifier

Feature vectors \mathbf{x} classified as $f(\mathbf{w}^\top \mathbf{x} + b) \in \{-1, +1\}$


f is a discontinuous step function

$$f(v) = \begin{cases} +1 & \text{if } v > 0 \\ -1 & \text{otherwise} \end{cases}$$

\mathbf{w}, b learned from training data

Multilayer Perceptrons

The Perceptron



Multilayer Perceptrons

The Perceptron – Limitations

Only two classes

Linear decision boundaries

Learning never converges for non-separable data

Multilayer Perceptrons

Two-Layer Architecture

Replace f with continuous nonlinearity (e.g. $\tanh(\cdot)$)

Introduce layer of M such “Perceptrons” (hidden units)

Hidden units connected to layer of K output units




Image from Bishop 2006

Multilayer Perceptrons

Two-Layer Architecture

Output of m th hidden unit is $z_m(\mathbf{x}) = f(\mathbf{w}_m^\top \mathbf{x})$

Bias b included in \mathbf{w} and \mathbf{x} , $w_0 = b$, $x_0 = 1$




Image from Bishop 2006

Multilayer Perceptrons

Two-Layer Architecture

Output of k th output unit is $y_k(\mathbf{z}) = g(\mathbf{w}_k^\top \mathbf{z})$

Choice of g depends on problem (regression, classification)




Image from Bishop 2006

Multilayer Perceptrons

Two-Layer Architecture

Both f and g are differentiable

- ▶ Learn \mathbf{w} using gradient descent
- ▶ Gradients evaluated via error backpropagation

The Pylearn2 Library

Machine learning library with focus on DL


Written in Python, but interaction mostly in YAML

Open-source: <https://github.com/lisa-lab/pylearn2>

The Pylearn2 Library

MLP Regression Example

Use pylearn2 to train a MLP for regression



The Pylearn2 Library

MLP Regression Example

```
model: !obj:pylearn2.models.mlp.MLP {  
  nvis: 1, # one input unit x  
  layers: [ # two layers  
    !obj:pylearn2.models.mlp.Tanh { # tanh activations for hidden units  
      dim: 3, # use M=3 hidden units  
      layer_name: 'hidden',  
      irange: 0.1  
    },  
    !obj:pylearn2.models.mlp.Linear { # linear output layer for regression  
      dim: 1, # one output unit, K=1  
      layer_name: 'out',  
      irange: 0.1  
    }  
  ]  
}
```

The Pylearn2 Library


MLP Regression Example

```
# dataset contains the (x,y) pairs from the previous figure
dataset: &train !pk1: 'mlp_data_regression.pkl',
# train using batch gradient descent
algorithm: !obj:pylearn2.training_algorithms.bgd.BGD {
    conjugate: 1,
    batch_size: 50,
    line_search_mode: 'exhaustive',
    termination_criterion: !obj:pylearn2.termination_criteria.EpochCounter {
        max_epochs: 100 # train for 100 epochs
    }
}
```

The Pylearn2 Library

MLP Regression Example

Full example: <https://github.com/cpra/cvsp-vo-slides>




Convolutional Neural Networks

Motivation

Above MLP architecture does not scale to images

- ▶ VGA resolution, $M = 1000$: ~ 300 million params




Convolutional Neural Networks

Motivation

Nearby pixels are closely correlated

Exploit topology through locally connected layers




Convolutional Neural Networks

Convolutional Layers

Convolutional layer consists of F feature maps

Each map consists of M hidden units with **shared** weights

Hidden units have a local receptive field



Convolutional Neural Networks

Convolutional Layers

Number of params independent of image resolution and M

$F = 100$, 15×15 receptive field: $\sim 23,000$ params

Feature map evaluation equals

- ▶ Convolution with kernel \mathbf{w}_f (hence the name)
- ▶ Followed by nonlinearity $f(\cdot) = \max(0, \cdot)$ (ReLU)




Image from Krizhevsky, Sutskever, and Hinton 2012


Convolutional Neural Networks

Pooling Layers

After convolutional layer

Hidden units pool information locally (e.g. max, avg)

Data reduction, robustness to small translations




Convolutional Neural Networks

Local Contrast Normalization Layers

Between convolution and pooling layers

Normalize responses spatially or over adjacent feature maps

Produce more expressive, robust features



Convolutional Neural Networks

Architecture

MLPs with above layers are Convolutional Neural Networks (CNNs)

Usually several Conv \Rightarrow Norm \Rightarrow Pooling blocks





Image adapted from Kavukcuoglu 2011

Convolutional Neural Networks

Architecture

Blocks followed by a traditional MLP



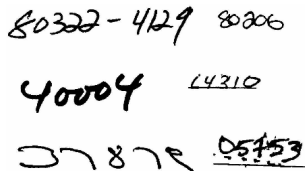
Deep Learning Applications

Zip Code Recognition

Zip code recognition from images

Among first applications of CNNs (1989!)

- ▶ DL and CNNs are “old” concepts
- ▶ Now successful due to available data & processing power



Handwritten zip codes from the MNIST dataset, showing examples of the original handwritten digits and the corresponding machine-recognized digits. The first row shows '80322-4129' and '80206'. The second row shows '40004' and '14310'. The third row shows '37879' and '05153'.

Image from LeCun et al. 1989

Deep Learning Applications

Zip Code Recognition




Image from LeCun et al. 1989

Deep Learning Applications

Zip Code Recognition

Let's try this ourselves

<http://cs.stanford.edu/people/karpathy/convnetjs/demo/mnist.html>



Deep Learning Applications

Face Recognition

Large market (expected \$6.5B by 2018)

- ▶ Security, law enforcement, HCI, ...

Complicated task

- ▶ Pose, occlusions, aging, expressions, accessories



Image adapted from <http://vis-www.cs.umass.edu/lfw/>

Deep Learning Applications

Face Recognition – Taigman et al. 2013

Face recognition via 3D face frontalization and DL

- ▶ Conv \Rightarrow Pooling \Rightarrow Conv for low-level features
- ▶ Locally connected layers for high-level features
- ▶ ~ 120 million parameters, ~ 4 million training images




Image from Taigman et al. 2013

Human-level face verification performance





Image from Taigman et al. 2013

Deep Learning Applications

Object Recognition

Object recognition on the ImageNet database

- ▶ ~ 14 million images categorized hierarchically
- ▶ Annual object recognition challenges



IMAGENET

Image from <http://web.eecs.umich.edu/~jiadeng/>

Deep Learning Applications

Object Recognition – Krizhevsky, Sutskever, and Hinton 2012

Object classification and localization via DL

Won the 2012 challenge by a large margin

~ 60 million parameters, ~ 1.2 million training images



Deep Learning Applications

Object Recognition – Krizhevsky, Sutskever, and Hinton 2012

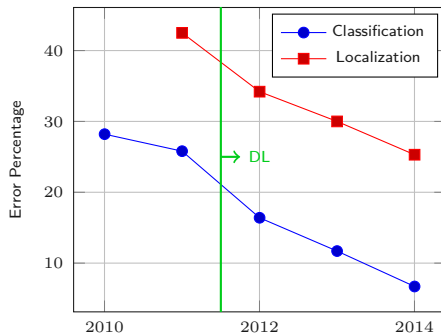
Implemented using cuda-convnet

- ▶ <https://code.google.com/p/cuda-convnet/>
- ▶ Open-source C++/CUDA implementation of CNNs
- ▶ Network structure definitions available on the webpage

Deep Learning Applications

Object Recognition

DL has lead to significant performance gains on ImageNet



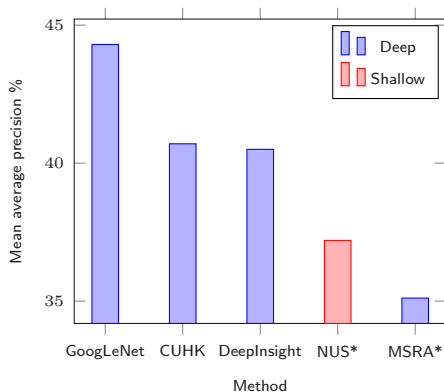
Online demo of 2013 winner of classification challenge

<http://www.clarifai.com/>

Deep Learning Applications

Object Recognition

Results of 2014 object detection challenge (excerpt)



Deep Learning Applications

Other Applications

Scene labeling

Action recognition

Speech recognition

...

Remarks

DL and CNNs are very powerful concepts

- ▶ State of the art results in many areas of application

Training CNNs is complex

- ▶ Some CNNs are trained on GPUs for weeks
- ▶ Large training datasets required
- ▶ Reason why DL became popular only recently

Several free implementations available

- ▶ pylearn2, convnetjs, cuda-convnet, caffe, torch7, ...

Bibliography I

Bengio, Yoshua, Ian Goodfellow, and Aaron Courville (2014).
Deep Learning (draft). MIT Press.

Bishop, Christopher (2006). **Pattern recognition and machine learning**. Springer.

Kavukcuoglu, Koray (2011). **Learning feature hierarchies for object recognition**. PhD thesis.

Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E Hinton (2012).
Imagenet classification with deep convolutional neural networks. NIPS.

LeCun, Yann et al. (1989). **Backpropagation applied to handwritten zip code recognition**. Neural computation.

Sermanet, Pierre et al. (2012). **Pedestrian Detection with Unsupervised Multi-Stage Feature Learning**. CoRR.

Taigman, Yaniv et al. (2013). **Deepface: Closing the gap to human-level performance in face verification**. CVPR.