

Computer Vision Systems Programming VO

Deep Learning

Christopher Pramerdorfer

Computer Vision Lab, Vienna University of Technology

Topics

Deep learning motivation



Multilayer perceptrons

Pylearn2 library

Convolutional neural networks

Deep learning applications


80322-4129
40004
3787e



Images from LeCun et al. 1989, Taigman et al. 2013, image-net.org

Object Recognition

Traditional Approach



Object Recognition

Traditional Approach

Problem: how to choose the representation/features?

“General” features not optimal

- ▶ Not tuned to task at hand, low-level

Designing task-specific features is complex

- ▶ Virtually impossible to do optimally

Object Recognition

Deep Learning

Solution: learn representation as well

Learning high-level representations directly is difficult

Deep Learning (DL) solves this

- ▶ By learning a hierarchy of representations
- ▶ Layers in hierarchy build upon each other

Object Recognition

Deep Learning

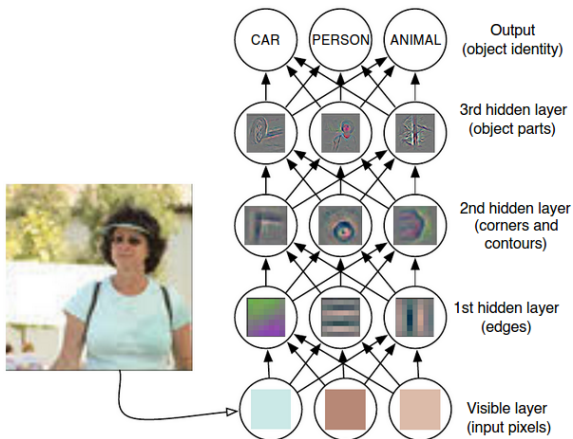


Image from Bengio, Goodfellow, and Courville 2014

Object Recognition

Deep Learning

n levels of features/representations

Learned jointly with the output mapping



Multilayer Perceptrons

DL is usually realized using MultiLayer Perceptrons (MLPs)




Image from Bishop 2006

Multilayer Perceptrons

The Perceptron

Binary linear classifier

Feature vectors \mathbf{x} classified as $f(\mathbf{w}^\top \mathbf{x} + b) \in \{-1, +1\}$


f is a discontinuous step function

$$f(v) = \begin{cases} +1 & \text{if } v > 0 \\ -1 & \text{otherwise} \end{cases}$$

\mathbf{w}, b learned from training data

Multilayer Perceptrons

The Perceptron



Multilayer Perceptrons

The Perceptron – Limitations

Only two classes

Linear decision boundaries

Learning never converges for non-separable data

Multilayer Perceptrons

Two-Layer Architecture

Replace f with continuous nonlinearity (e.g. $\tanh(\cdot)$)

Introduce layer of M such “Perceptrons” (hidden units)

Hidden units connected to layer of K output units




Image from Bishop 2006

Multilayer Perceptrons

Two-Layer Architecture

Output of m th hidden unit is $z_m(\mathbf{x}) = f(\mathbf{w}_m^\top \mathbf{x})$

Bias b included in \mathbf{w} and \mathbf{x} , $w_0 = b$, $x_0 = 1$




Image from Bishop 2006

Multilayer Perceptrons

Two-Layer Architecture

Output of k th output unit is $y_k(\mathbf{z}) = g(\mathbf{w}_k^\top \mathbf{z})$

Choice of g depends on problem (regression, classification)




Image from Bishop 2006

Multilayer Perceptrons

Two-Layer Architecture

Both f and g are differentiable

- ▶ Learn \mathbf{w} using gradient descent
- ▶ Gradients evaluated via error backpropagation

The Pylearn2 Library

Machine learning library with focus on DL


Written in Python, but interaction mostly in YAML

Open-source: <https://github.com/lisa-lab/pylearn2>

The Pylearn2 Library

MLP Regression Example

We will use pylearn2 to train a MLP for regression



The Pylearn2 Library

MLP Regression Example

```
model: !obj:pylearn2.models.mlp.MLP {  
  nvis: 1, # one input unit x  
  layers: [ # two layers  
    !obj:pylearn2.models.mlp.Tanh { # tanh activations for hidden units  
      dim: 3, # use M=3 hidden units  
      layer_name: 'hidden',  
      irange: 1  
    },  
    !obj:pylearn2.models.mlp.Linear { # linear output layer for regression  
      dim: 1, # one output unit, K=1  
      layer_name: 'out',  
      irange: 1  
    }  
  ]  
}
```

The Pylearn2 Library


MLP Regression Example

```
# dataset contains the (x,y) pairs from the previous figure
dataset: &train !pk1: 'mlp_data_regression.pkl',
# train using batch gradient descent
algorithm: !obj:pylearn2.training_algorithms.bgd.BGD {
    conjugate: 1,
    batch_size: 50,
    line_search_mode: 'exhaustive',
    termination_criterion: !obj:pylearn2.termination_criteria.EpochCounter {
        max_epochs: 100 # train for 100 epochs
    }
}
```

The Pylearn2 Library

MLP Regression Example

Full example: <https://github.com/cpra/cvsp-vo-slides>



Bibliography I

Bengio, Yoshua, Ian Goodfellow, and Aaron Courville (2014).

Deep Learning (draft). MIT Press.

Bishop, Christopher (2006). **Pattern recognition and machine learning.** Springer.

LeCun, Yann et al. (1989). **Backpropagation applied to handwritten zip code recognition.** Neural computation.

Taigman, Yaniv et al. (2013). **Deepface: Closing the gap to human-level performance in face verification.** CVPR.