

# Computer Vision Systems Programming VO

## Object Recognition – Deep Learning

Christopher Pramerdorfer

Computer Vision Lab, Vienna University of Technology

# Topics



Deep learning motivation

Multilayer perceptrons

Convolutional neural networks

Deep learning applications


80322-4129  
40004  
07879



Images from LeCun et al. 1989, Taigman et al. 2013, image-net.org

# Object Recognition

## Traditional Approach



# Object Recognition

## Traditional Approach

Problem: how to choose the representation/features?

“General” features not optimal

- ▶ Not tuned to task at hand, low-level

Designing task-specific features is complex

- ▶ Virtually impossible to do optimally

Solution: learn representation as well

Learning high-level representations directly is difficult

**Deep Learning** (DL) solves this

- ▶ By learning a hierarchy of representations
- ▶ Layers in hierarchy build upon each other

# Object Recognition

## Deep Learning

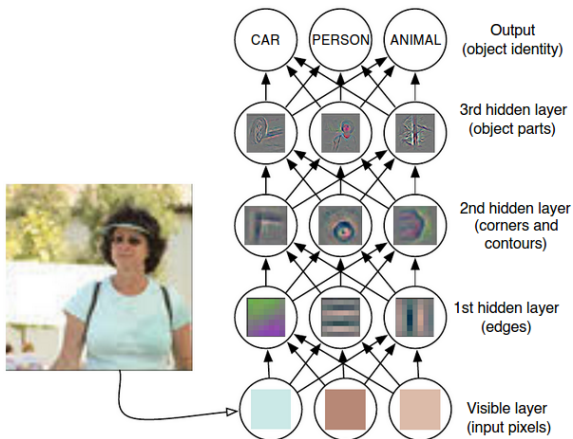


Image from Bengio, Goodfellow, and Courville 2014

# Object Recognition

## Deep Learning

$n$  levels of features/representations

Learned jointly with the output mapping



Methods based on DL are state of the art

- ▶ For almost all recognition-related CV tasks

“From now on, DL has to be considered as the primary candidate in essentially any visual recognition task”  
[Razavian et al. 2014]



# Multilayer Perceptrons

DL is realized using MultiLayer Perceptrons (MLPs)




Image from Bishop 2006

# Multilayer Perceptrons

## The Perceptron

Binary linear classifier

Feature vectors  $\mathbf{x}$  classified as  $f(\mathbf{w}^\top \mathbf{x} + b) \in \{-1, +1\}$


$f$  is a discontinuous step function

$$f(v) = \begin{cases} +1 & \text{if } v > 0 \\ -1 & \text{otherwise} \end{cases}$$

$\mathbf{w}, b$  learned from training data

# Multilayer Perceptrons

## The Perceptron



# Multilayer Perceptrons

## The Perceptron – Limitations

Only two classes

Linear decision boundaries

Learning never converges for non-separable data

# Multilayer Perceptrons

## Two-Layer Architecture

Replace  $f$  with continuous nonlinearity (e.g.  $\tanh(\cdot)$ )

Introduce layer of  $M$  such “Perceptrons” (hidden units)

Hidden units connected to layer of  $K$  output units




Image from Bishop 2006

# Multilayer Perceptrons

## Two-Layer Architecture

Output of  $m$ th hidden unit is  $z_m(\mathbf{x}) = f(\mathbf{w}_m^\top \mathbf{x})$

- Bias  $b$  included in  $\mathbf{w}$  and  $\mathbf{x}$ ,  $w_0 = b$ ,  $x_0 = 1$




Image from Bishop 2006

# Multilayer Perceptrons

## Two-Layer Architecture

Output of  $k$ th output unit is  $y_k(\mathbf{z}) = g(\mathbf{w}_k^\top \mathbf{z})$

Choice of  $g$  depends on problem (regression, classification)




Image from Bishop 2006

# Multilayer Perceptrons

## Two-Layer Architecture

Both  $f$  and  $g$  are differentiable

- ▶ Learn  $\mathbf{w}$  using gradient descent
- ▶ Gradients evaluated via error backpropagation



# The Pylearn2 Library

Machine learning library with focus on DL


Written in Python, but interaction mostly in YAML

Open-source: <https://github.com/lisa-lab/pylearn2>

# The Pylearn2 Library

## MLP Regression Example

Use pylearn2 to train a MLP for regression



# The Pylearn2 Library

## MLP Regression Example

```
model: !obj:pylearn2.models.mlp.MLP {  
  nvis: 1, # one input unit x  
  layers: [ # two layers  
    !obj:pylearn2.models.mlp.Tanh { # tanh activations for hidden units  
      dim: 3, # use M=3 hidden units  
      layer_name: 'hidden',  
      irange: 0.1  
    },  
    !obj:pylearn2.models.mlp.Linear { # linear output layer for regression  
      dim: 1, # one output unit, K=1  
      layer_name: 'out',  
      irange: 0.1  
    }  
  ]  
}
```

# The Pylearn2 Library


## MLP Regression Example

```
# dataset contains the (x,y) pairs from the previous figure
dataset: &train !pk1: 'mlp_data_regression.pkl',
# train using batch gradient descent
algorithm: !obj:pylearn2.training_algorithms.bgd.BGD {
    conjugate: 1,
    batch_size: 50,
    line_search_mode: 'exhaustive',
    termination_criterion: !obj:pylearn2.termination_criteria.EpochCounter {
        max_epochs: 100 # train for 100 epochs
    }
}
```

# The Pylearn2 Library

## MLP Regression Example

Full example: <https://github.com/cpra/cvsp-vo-slides>




# Convolutional Neural Networks

## Motivation

Above MLP architecture does not scale to images

- ▶ VGA resolution,  $M = 1000$ :  $\sim 300$  million params




# Convolutional Neural Networks

## Motivation

Nearby pixels are closely correlated

Exploit topology through locally connected layers



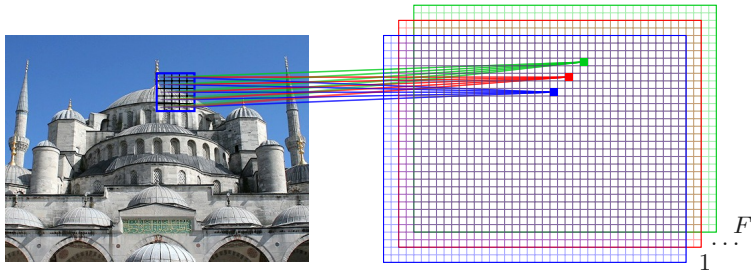
# Convolutional Neural Networks

## Convolutional Layers

Convolutional layer consists of  $F$  feature maps

Each map consists of  $M$  hidden units with **shared** weights

Hidden units have a local receptive field





# Convolutional Neural Networks

## Convolutional Layers

Number of params independent of image resolution and  $M$

$F = 100$ ,  $15 \times 15$  receptive field:  $\sim 23,000$  params

Feature map evaluation equals

- ▶ Convolution with kernel  $\mathbf{w}_f$  (hence the name)
- ▶ Followed by nonlinearity  $f(\cdot) = \max(0, \cdot)$  (ReLU)



Image from Krizhevsky, Sutskever, and Hinton 2012


# Convolutional Neural Networks

## Pooling Layers

After convolutional layer

Hidden units pool information locally (e.g. max, avg)

Data reduction, robustness to small translations




# Convolutional Neural Networks

## Local Contrast Normalization Layers

Between convolution and pooling layers

Normalize responses spatially or over adjacent feature maps

Produce more expressive, robust features



# Convolutional Neural Networks

## Architecture

MLPs with above layers are Convolutional Neural Networks (CNNs)

Usually several Conv  $\Rightarrow$  Norm  $\Rightarrow$  Pooling blocks

Later Conv layers operate on multiple feature maps


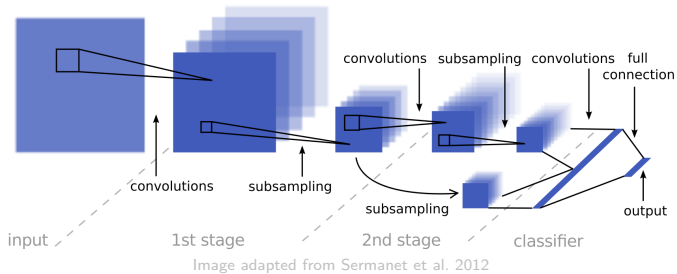


Image adapted from Kavukcuoglu 2011

# Convolutional Neural Networks

## Architecture

Blocks followed by a traditional MLP



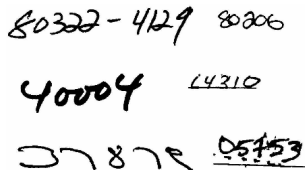
# Deep Learning Applications

## Zip Code Recognition

Zip code recognition from images

Among first applications of CNNs (1989!)

- ▶ DL and CNNs are “old” concepts
- ▶ Now successful due to available data & processing power



Handwritten zip codes from the MNIST dataset, showing examples of the original handwritten digits and the corresponding machine-predicted digits. The first row shows '80322-4129' and '80206'. The second row shows '40004' and '14310'. The third row shows '37879' and '05153'.

Image from LeCun et al. 1989

# Deep Learning Applications

## Zip Code Recognition




Image from LeCun et al. 1989

# Deep Learning Applications

## Zip Code Recognition

Let's try this ourselves

<http://cs.stanford.edu/people/karpathy/convnetjs/demo/mnist.html>





# Deep Learning Applications

## Face Recognition

Large market (expected \$6.5B by 2018)

- ▶ Security, law enforcement, HCI, ...

Complicated task

- ▶ Pose, occlusions, aging, expressions, accessories



Image adapted from <http://vis-www.cs.umass.edu/lfw/>

# Deep Learning Applications

Face Recognition – Taigman et al. 2013

Face recognition via 3D face frontalization and DL

- ▶ Conv  $\Rightarrow$  Pooling  $\Rightarrow$  Conv for low-level features
- ▶ Locally connected layers for high-level features
- ▶  $\sim 120$  million parameters,  $\sim 4$  million training images




Image from Taigman et al. 2013

### Human-level face verification performance





Image from Taigman et al. 2013

# Deep Learning Applications

## Object Recognition

### Object recognition on the ImageNet database

- ▶ ~ 14 million images categorized hierarchically
- ▶ Annual object recognition challenges




IMAGENET

Image from <http://web.eecs.umich.edu/~jiadeng/>

# Deep Learning Applications

## Object Recognition


DL has lead to significant performance gains on ImageNet



# Deep Learning Applications

## Object Recognition

Results of 2014 ImageNet object detection challenge (excerpt)



# Deep Learning Applications

Object Recognition – Krizhevsky, Sutskever, and Hinton 2012

Object classification and localization via DL

Won the 2012 challenge by a large margin

~ 60 million parameters, ~ 1.2 million training images



# Deep Learning Applications

Object Recognition – Krizhevsky, Sutskever, and Hinton 2012

Implemented using cuda-convnet

- ▶ <https://code.google.com/p/cuda-convnet/>
- ▶ Open-source C++/CUDA implementation of CNNs
- ▶ Network structure definitions available on the webpage



Online demo of 2013 winner of classification challenge

<http://www.clarifai.com/>

# Deep Learning Applications

## Other Applications

Scene labeling

Action recognition

Speech recognition

...

# Remarks

DL and CNNs are very powerful concepts

- ▶ State of the art results in many areas of application

Training CNNs is complex

- ▶ Some CNNs are trained on GPUs for weeks
- ▶ Large training datasets required
- ▶ Reason why DL became popular only recently

Several free implementations available

- ▶ pylearn2, convnetjs, cuda-convnet, caffe, torch7, ...

# Bibliography I

Bengio, Yoshua, Ian Goodfellow, and Aaron Courville (2014).

**Deep Learning (draft).** MIT Press.

Bishop, Christopher (2006). **Pattern recognition and machine learning.** Springer.

Kavukcuoglu, Koray (2011). **Learning feature hierarchies for object recognition.** PhD thesis.

Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E Hinton (2012). **Imagenet classification with deep convolutional neural networks.**

# Bibliography II

LeCun, Yann et al. (1989). **Backpropagation applied to handwritten zip code recognition.**

Razavian, Ali Sharif et al. (2014). **CNN Features off-the-shelf: an Astounding Baseline for Recognition.**

Sermanet, Pierre et al. (2012). **Pedestrian Detection with Unsupervised Multi-Stage Feature Learning.**

Taigman, Yaniv et al. (2013). **Deepface: Closing the gap to human-level performance in face verification.**